

NDC { Oslo }

# » Event Sourcing with Azure Cosmos DB

Sander Molenkamp  
@amolenk



0 2 3 7 1,-

LANDIS & GYR

kWh

CL2f3

No. 35042469



P.G.E.M. TYPE W № 187249



Kamstrup

000 143 1 kWh  
1 L2 L3 T2

ZABF001586645911

KAMSTRUP 382JxC

684-31A-J2-C3-081

Sn: 15866459

1-2-3x230/400V

0.25-5(105)A

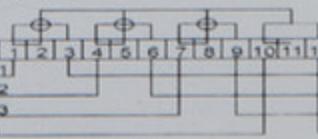
LED 1000 imp/kWh

Klasse A

2011W40

50/60 Hz

EN 50470-1/3



-40 - +55°C

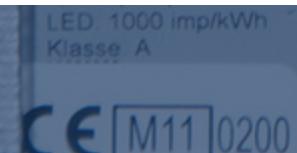
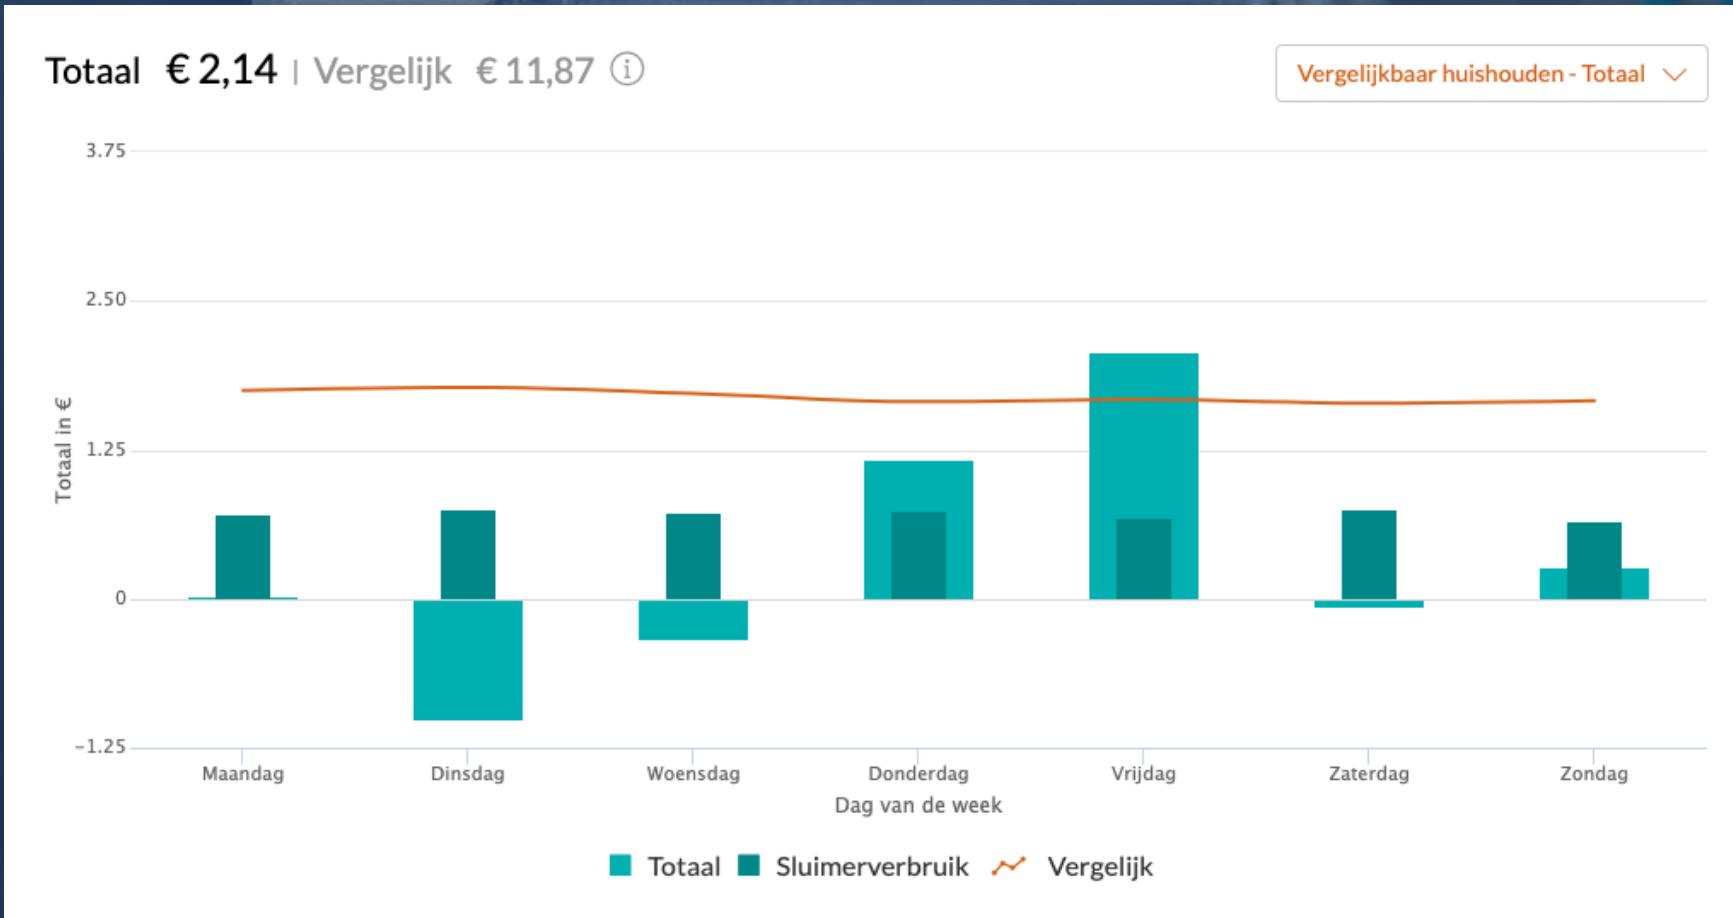
T10063

dlms



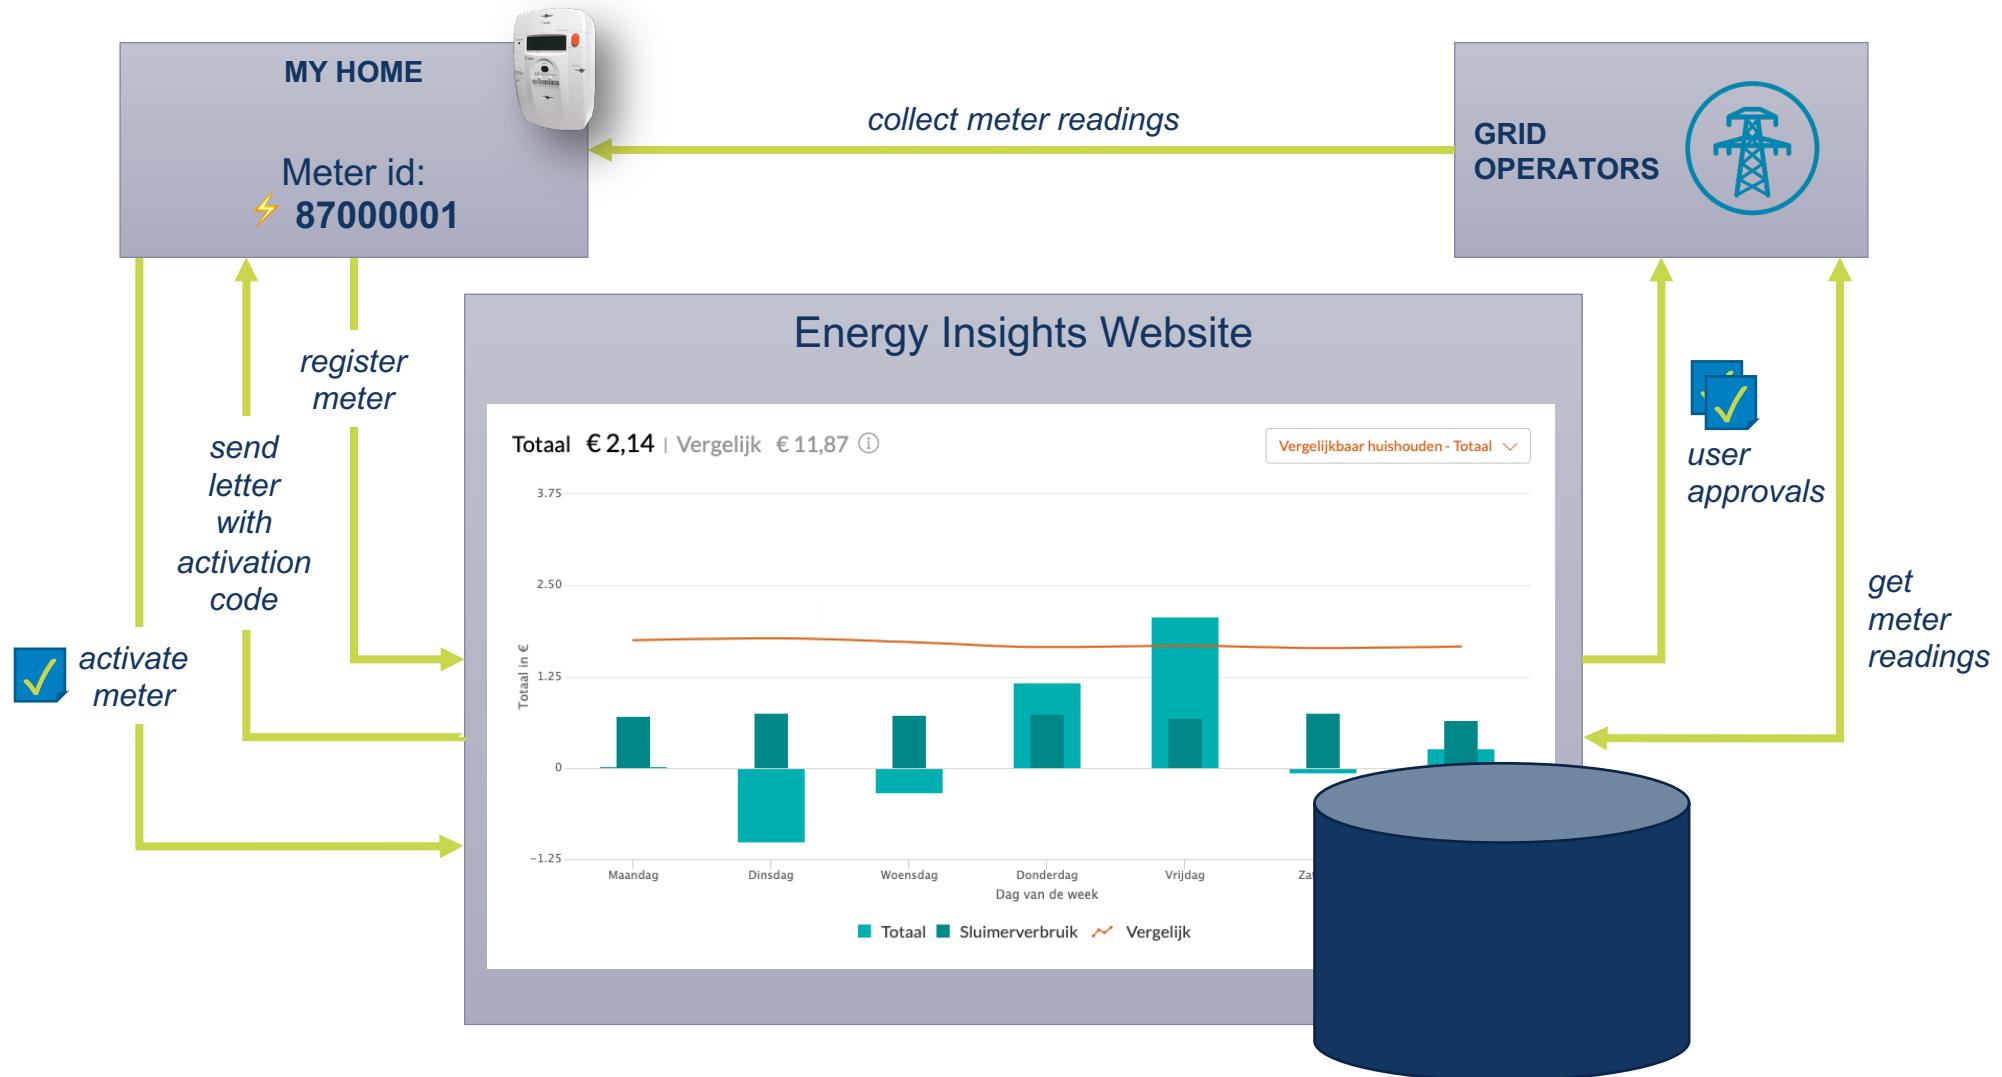
CE M11 0200

# Energy insights

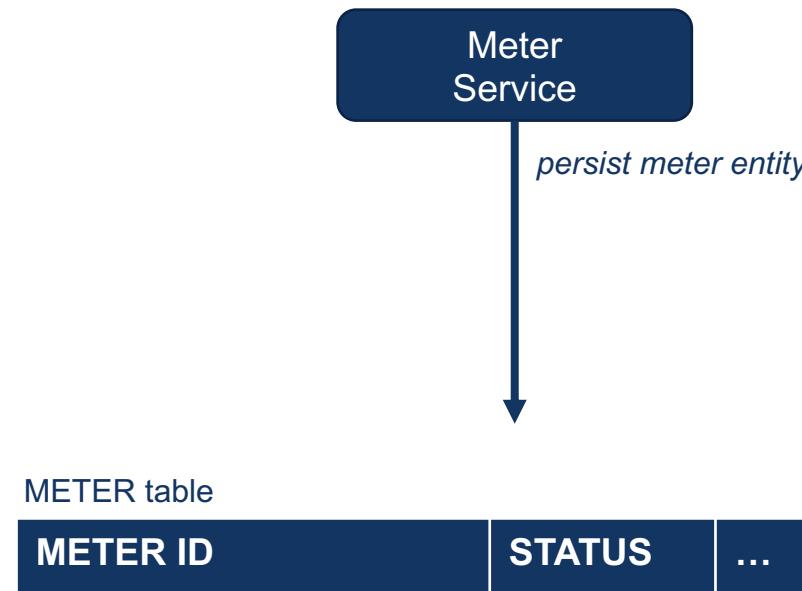


-40 - +55°C  
T10063 adlms

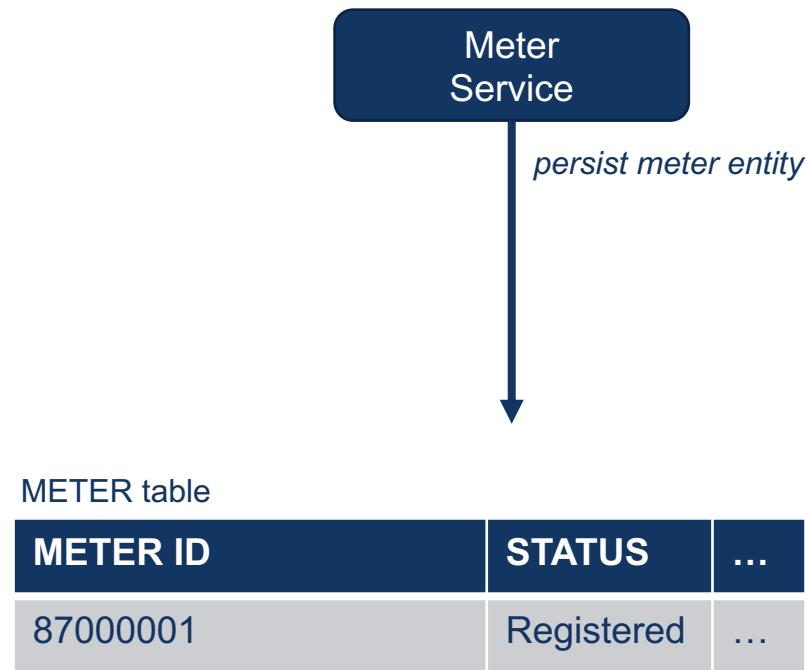
# Energy insights



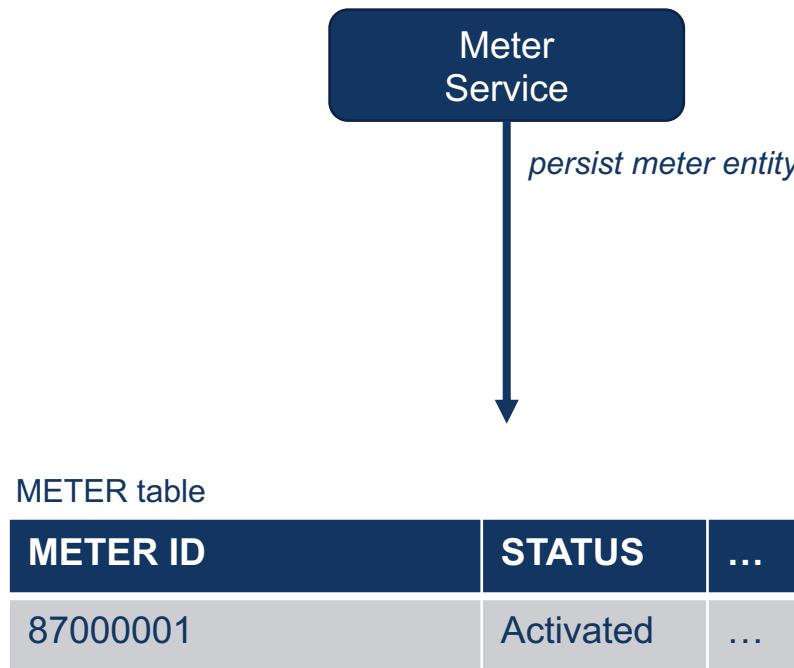
# ▲ Saving current state



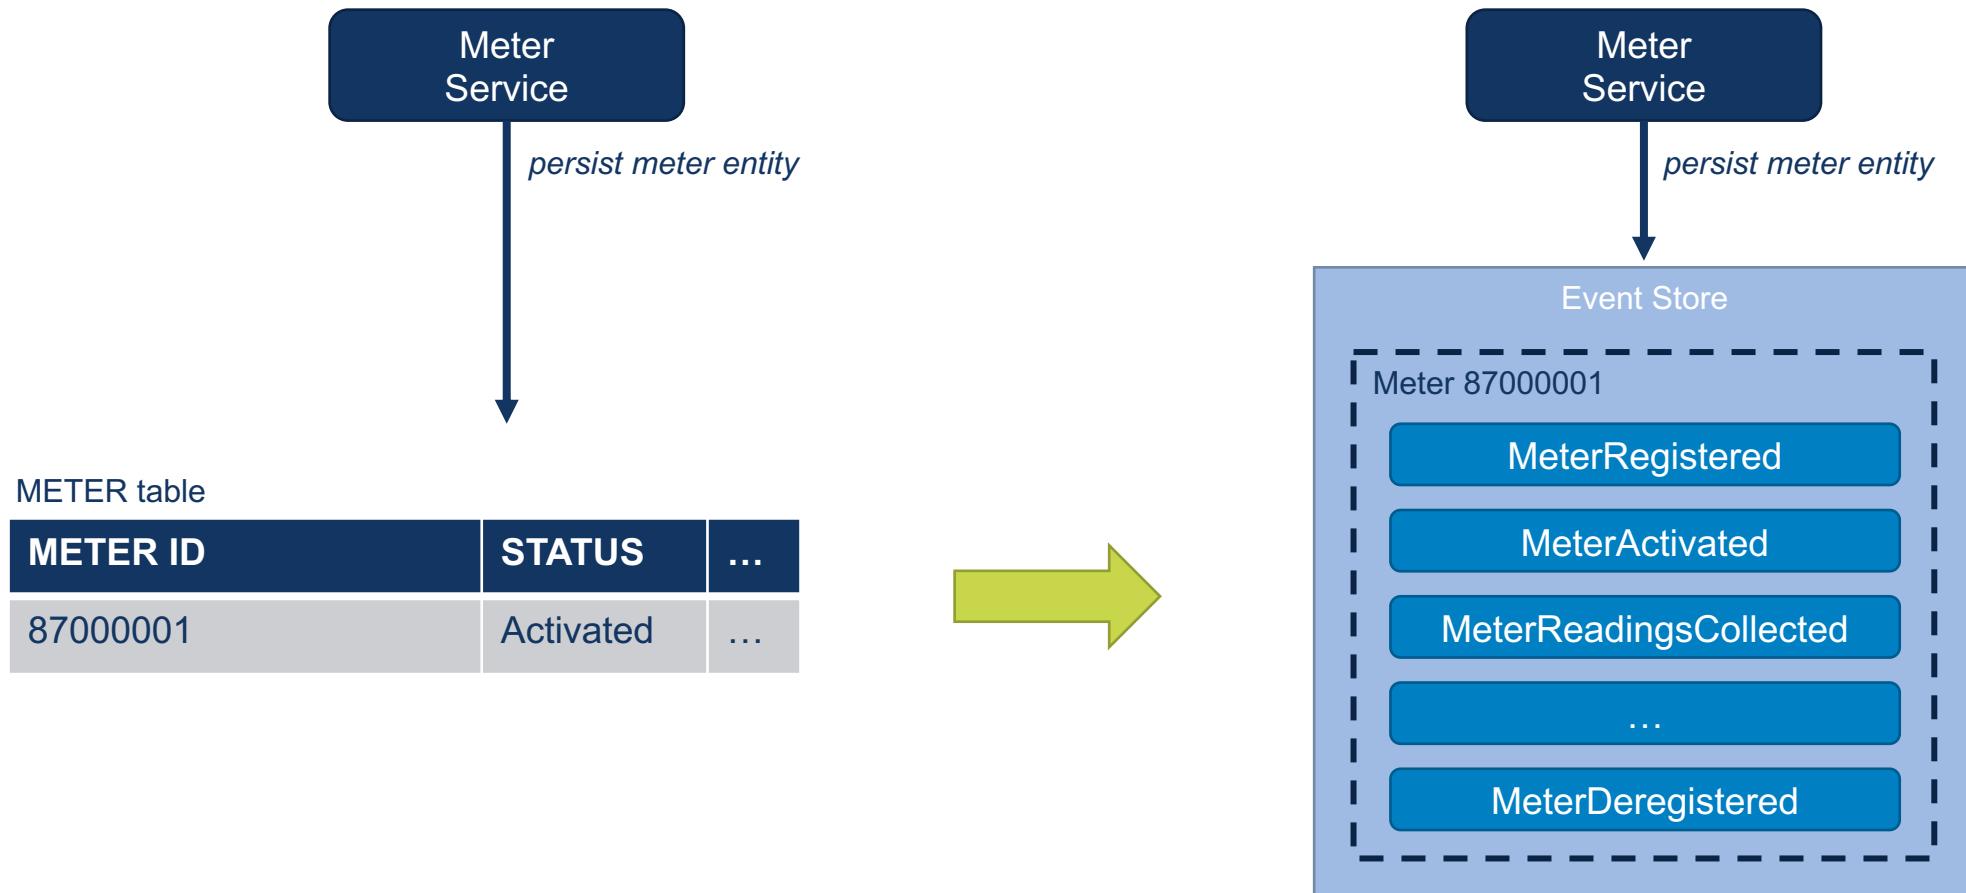
# ▲ Saving current state



# ▲ Saving current state



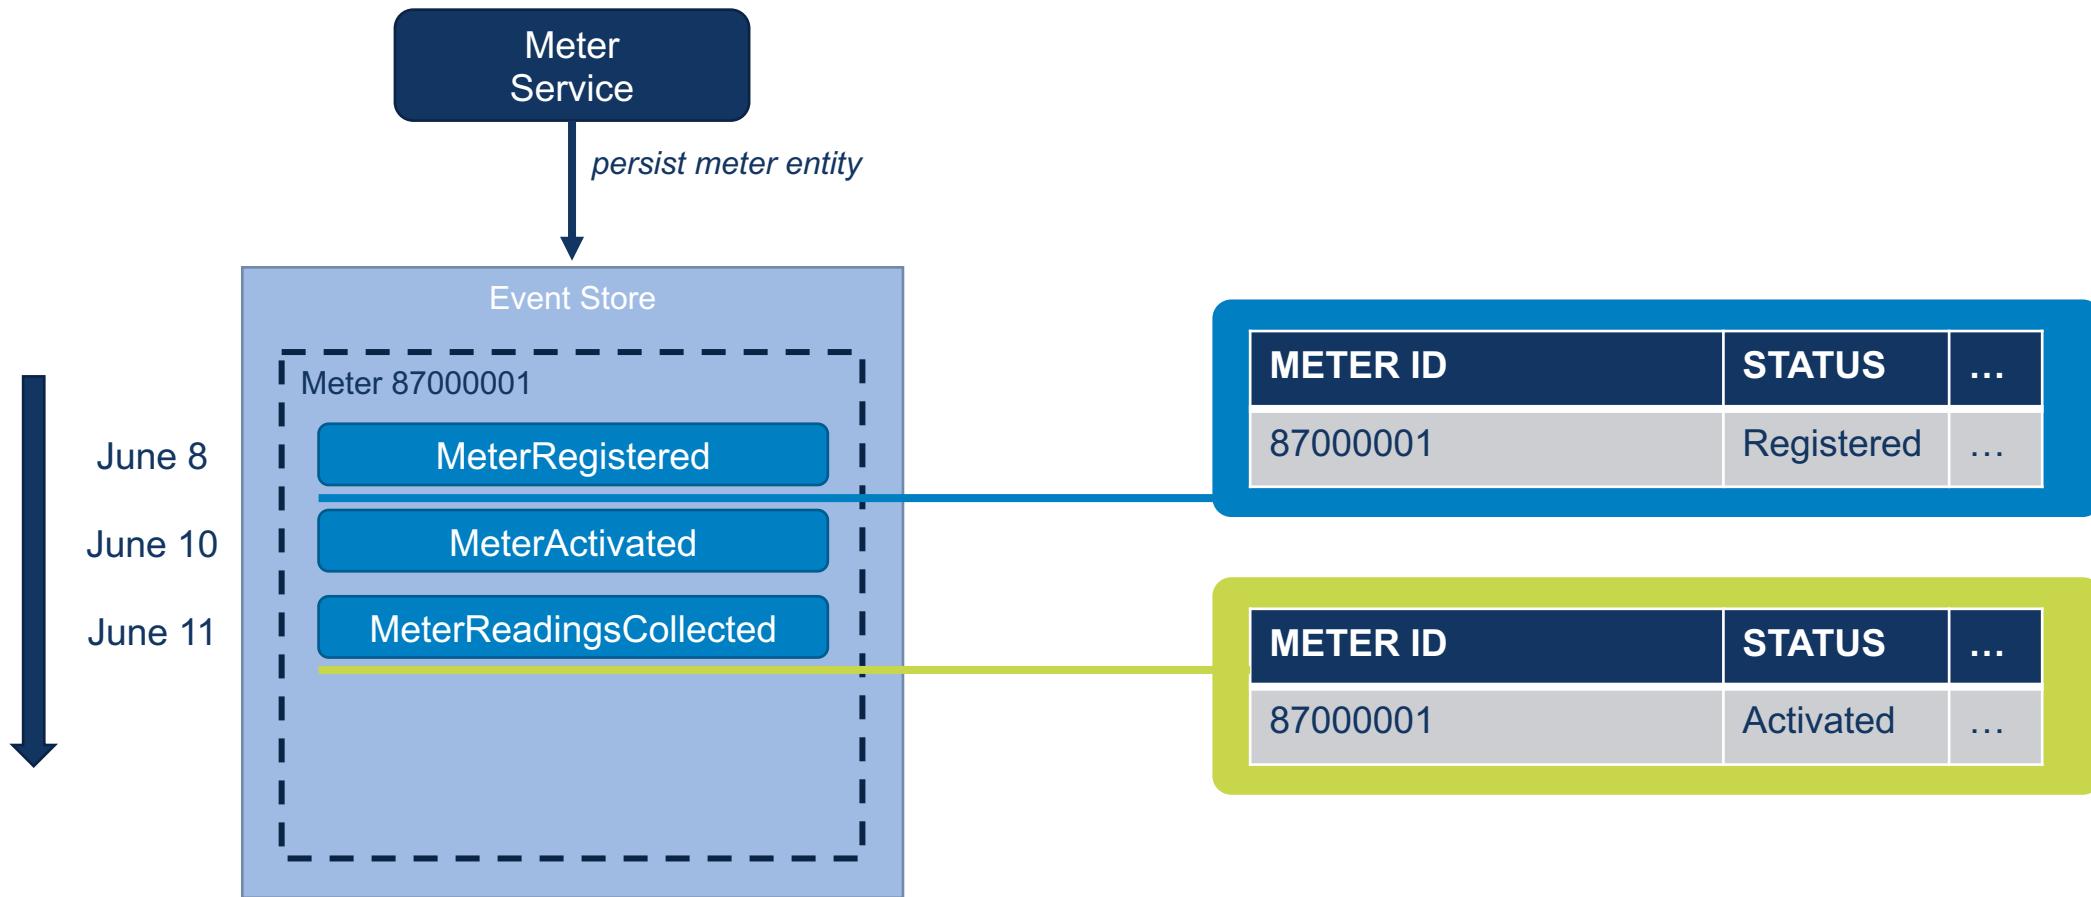
# ► Storing state as a sequence of events



## ◀ MeterRegistered event

```
{  
  "MeterId": "87000001",  
  "PostalCode": "1000 AA",  
  "HouseNumber": "25",  
  "ActivationCode": "542-484",  
  "Timestamp": "2020-04-17T05:37:57.775153Z"  
}
```

# ► Replay events to get current state



## ► Event sourcing

- Immutable stream of events
  - Just like an accountant you never change past events
- 100% accurate audit logging
- Easy temporal queries
  - Reconstructing historical state is easy by replaying a part of a stream

# Let's build an Event Store!

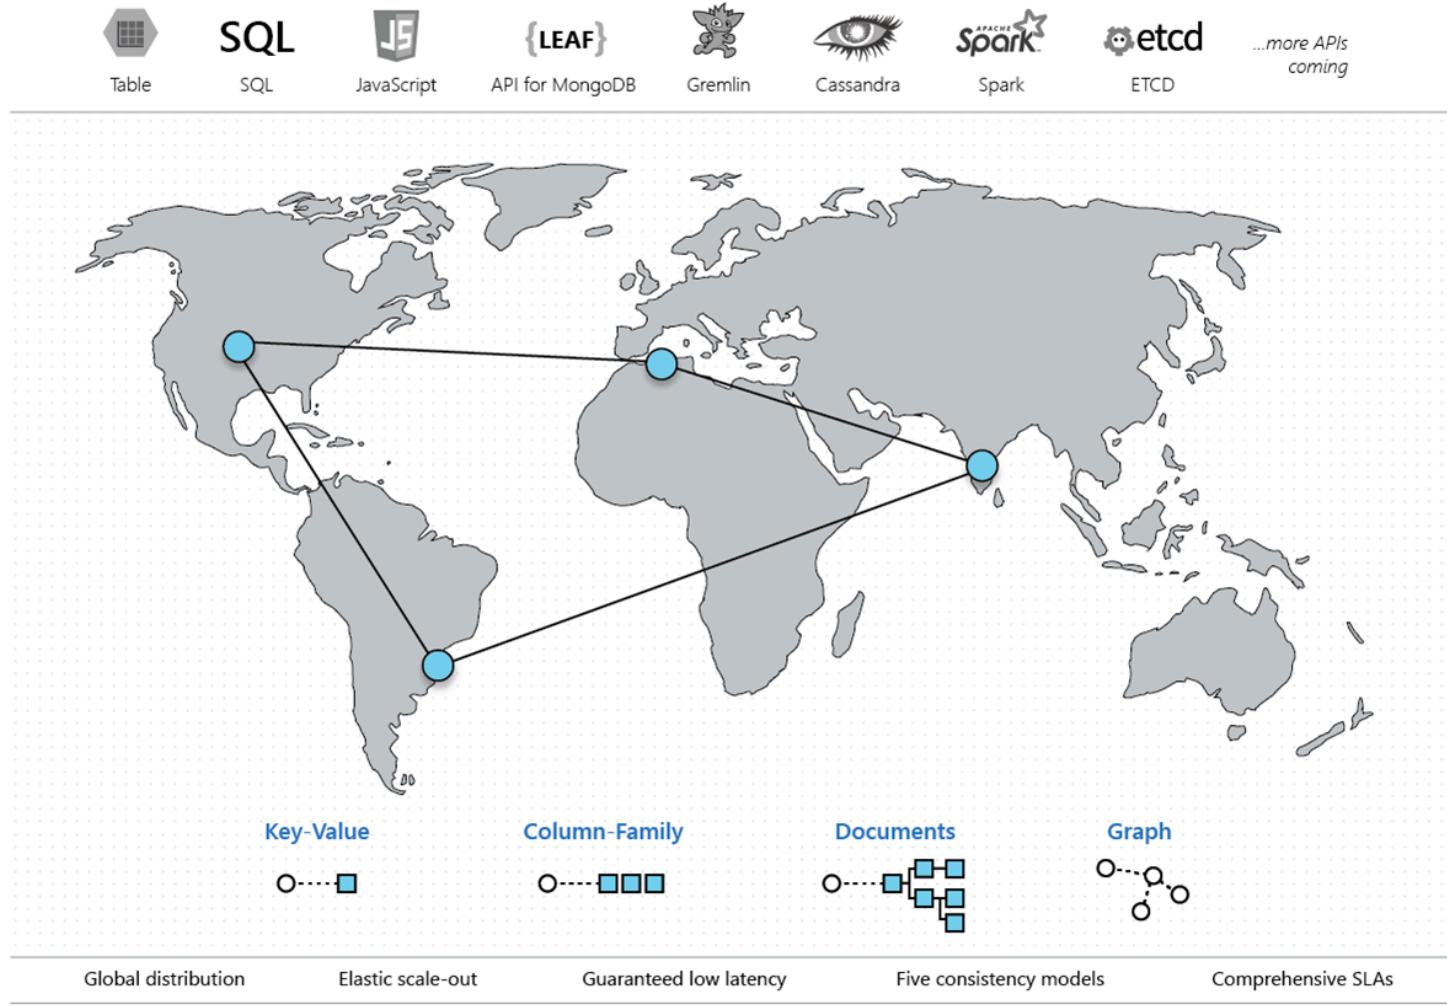


Building your own event store is a hands-on way to increase your understanding of the core concepts.

```
public interface IEventStore
{
    Task<EventStream> LoadStreamAsync(string streamId);

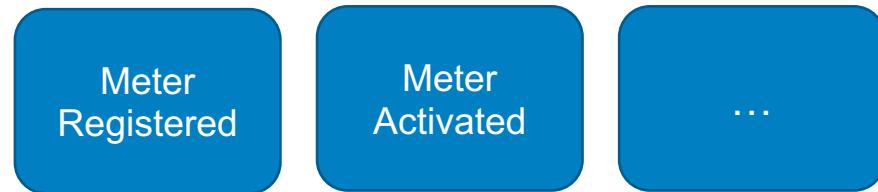
    Task AppendToStreamAsync(
        string streamId,
        int expectedVersion,
        IEnumerable<IEvent> events);
}
```

# Azure Cosmos DB



# Event sourcing using Cosmos DB

- We will use the SQL API
  - Each JSON item in the database represents a single event



- C# / SDK v3
- Each JSON item contains additional metadata
  - Stream ID
  - Stream version
  - Event type/name

# Choosing a partition key

Add Container X

**Start at \$24/mo per database, multiple containers included** [More details](#)

\* Database id ①  
 Create new  Use existing  
Type a new database id

Provision database throughput ①

\* Throughput (400 - 1,000,000 RU/s) ①  
 Autopilot (preview)  Manual  
400

Estimated spend (USD): **\$0.032 hourly / \$0.77 daily / \$23.04 monthly** (1 region, 400RU/s, \$0.00008/RU)

\* Container id ①  
e.g., Container1

\* Partition key ①  
e.g., /address/zipCode  
 My partition key is larger than 100 bytes

Unique keys ①  
+ Add unique key

## AppendToStream( streamId, expectedVersion, events)

### TRANSACTION SCOPE

1. Get current stream version
2. Verify that version matched expected version
3. If so, write events to store

- Use Stored Procedure for transaction

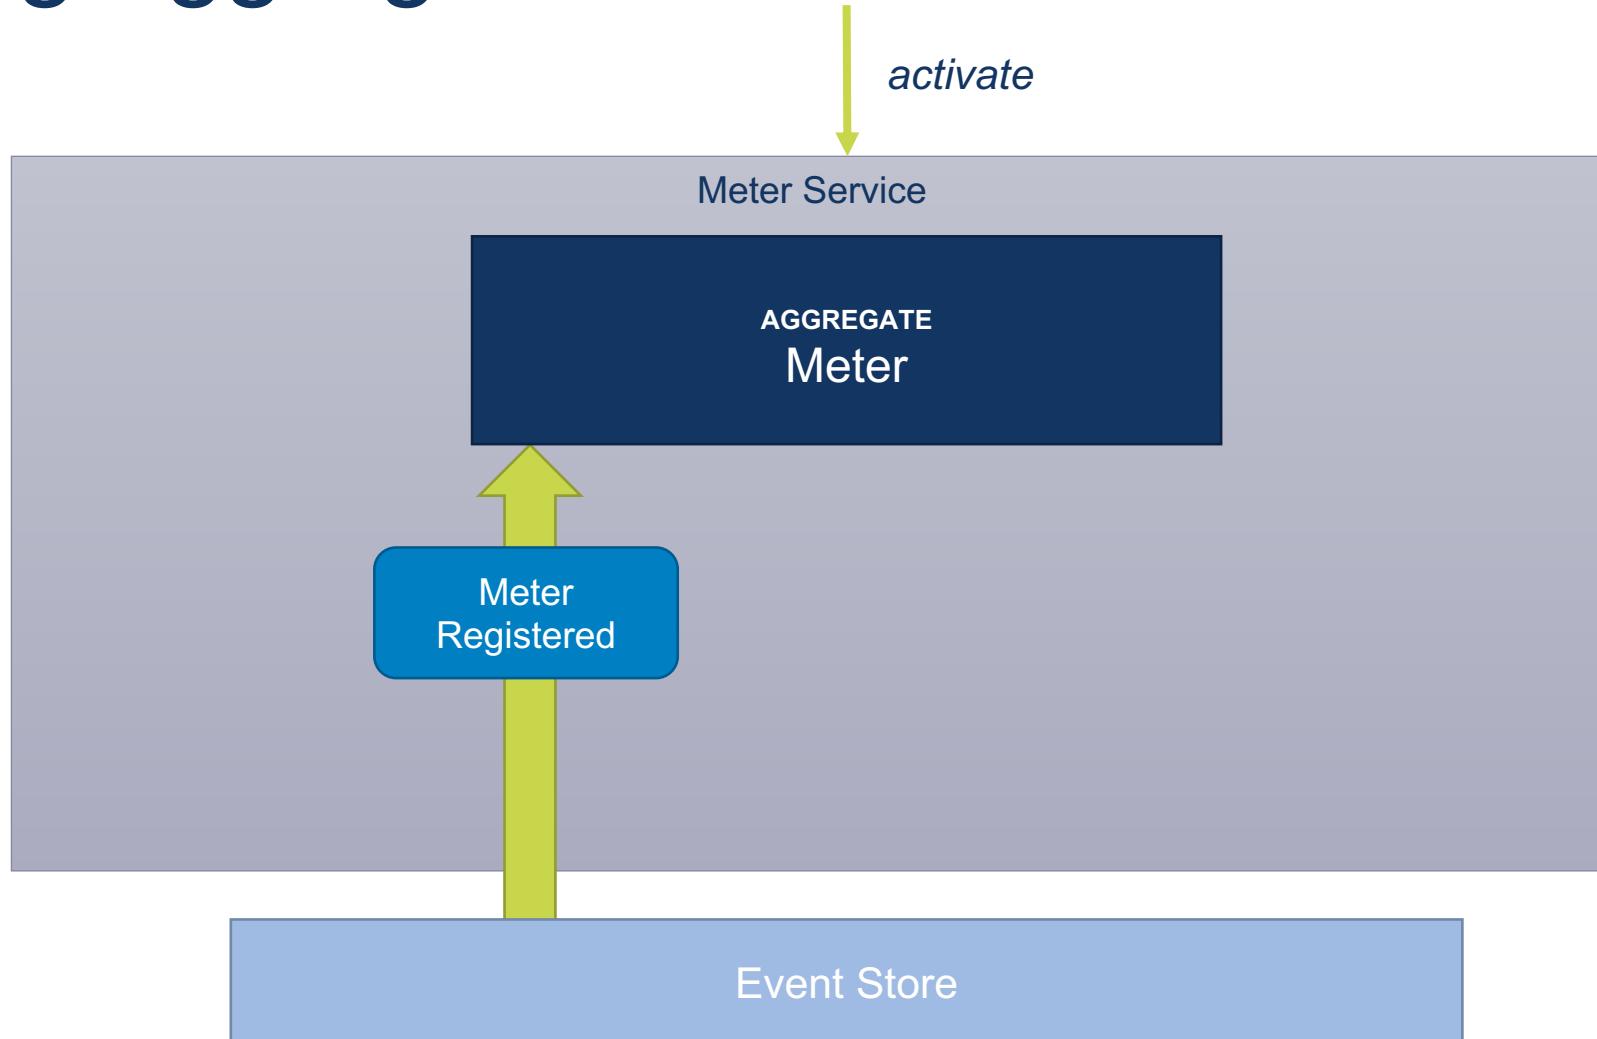
>>

# Writing & reading event streams

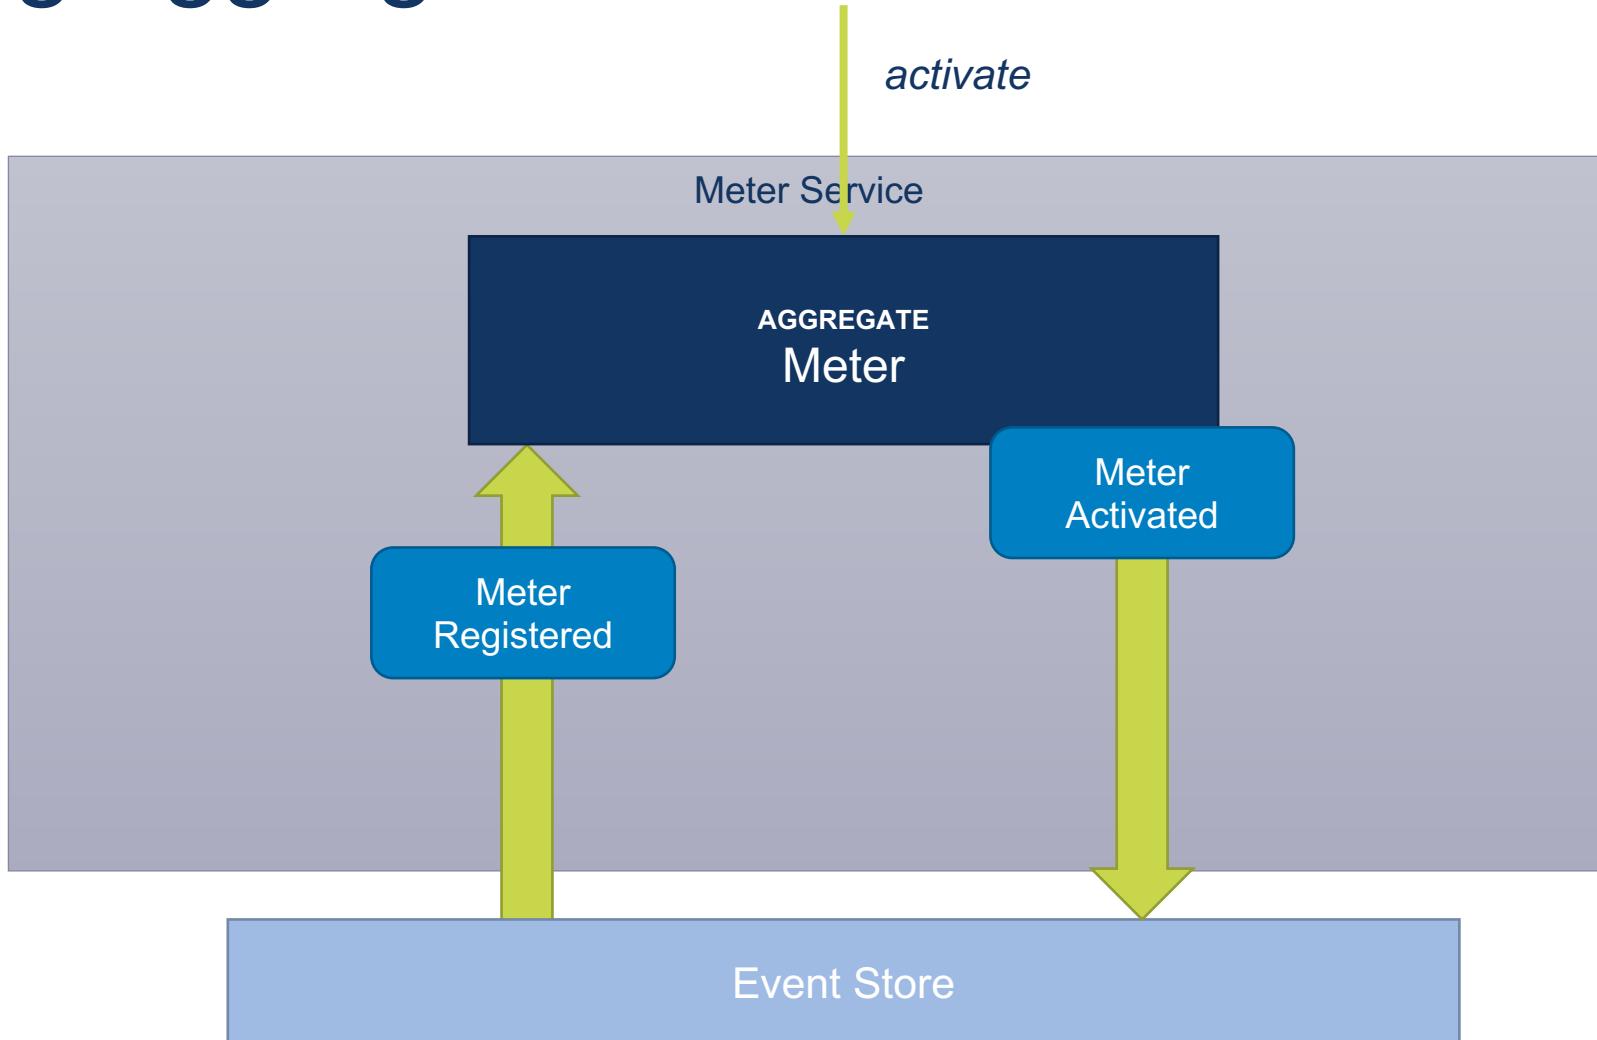
COSMOS DB - DEMO



# Using aggregates



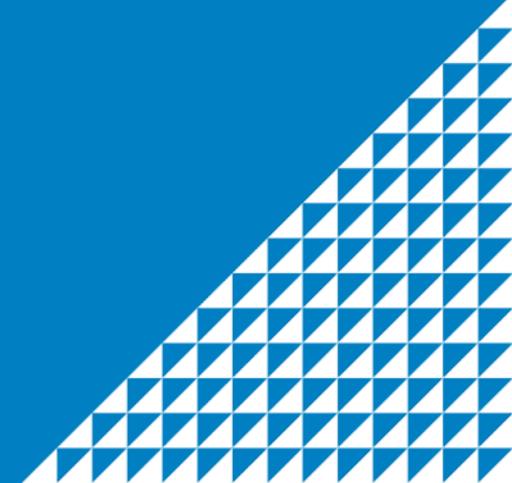
# Using aggregates



>>

# Adding business logic

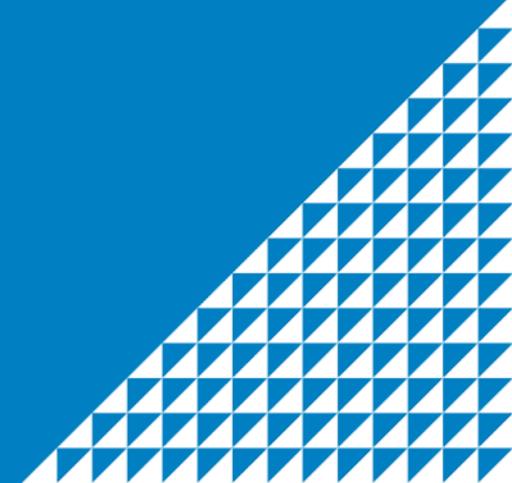
COSMOS DB - DEMO



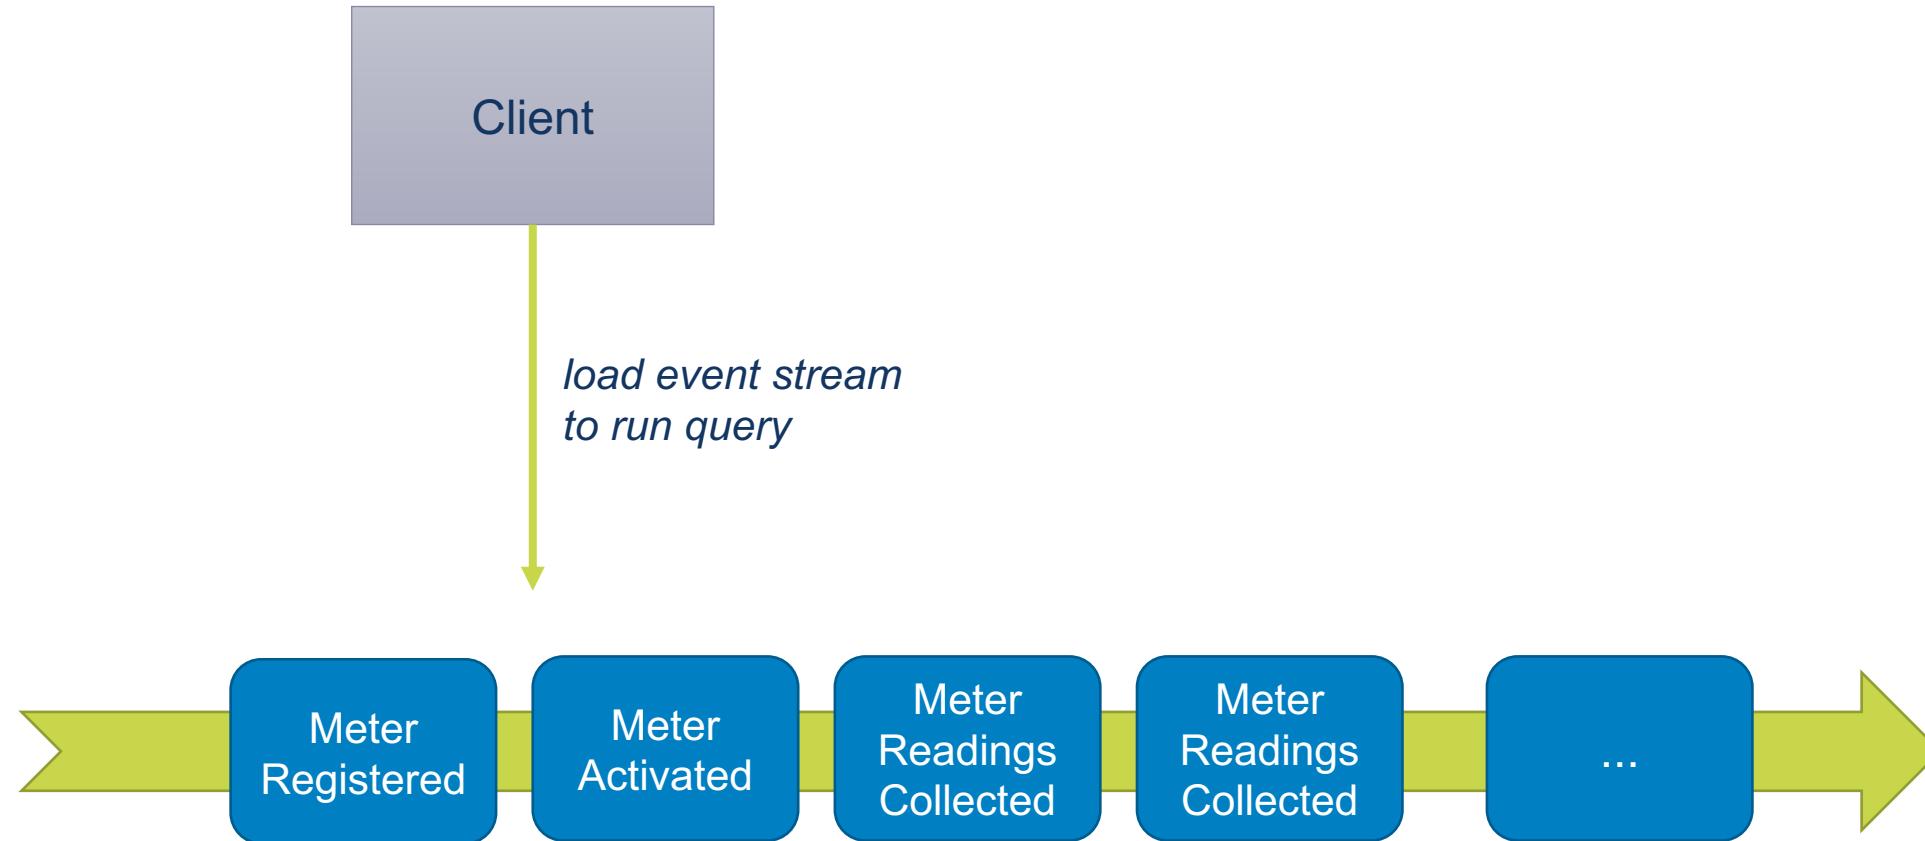
>>

# Querying event streams

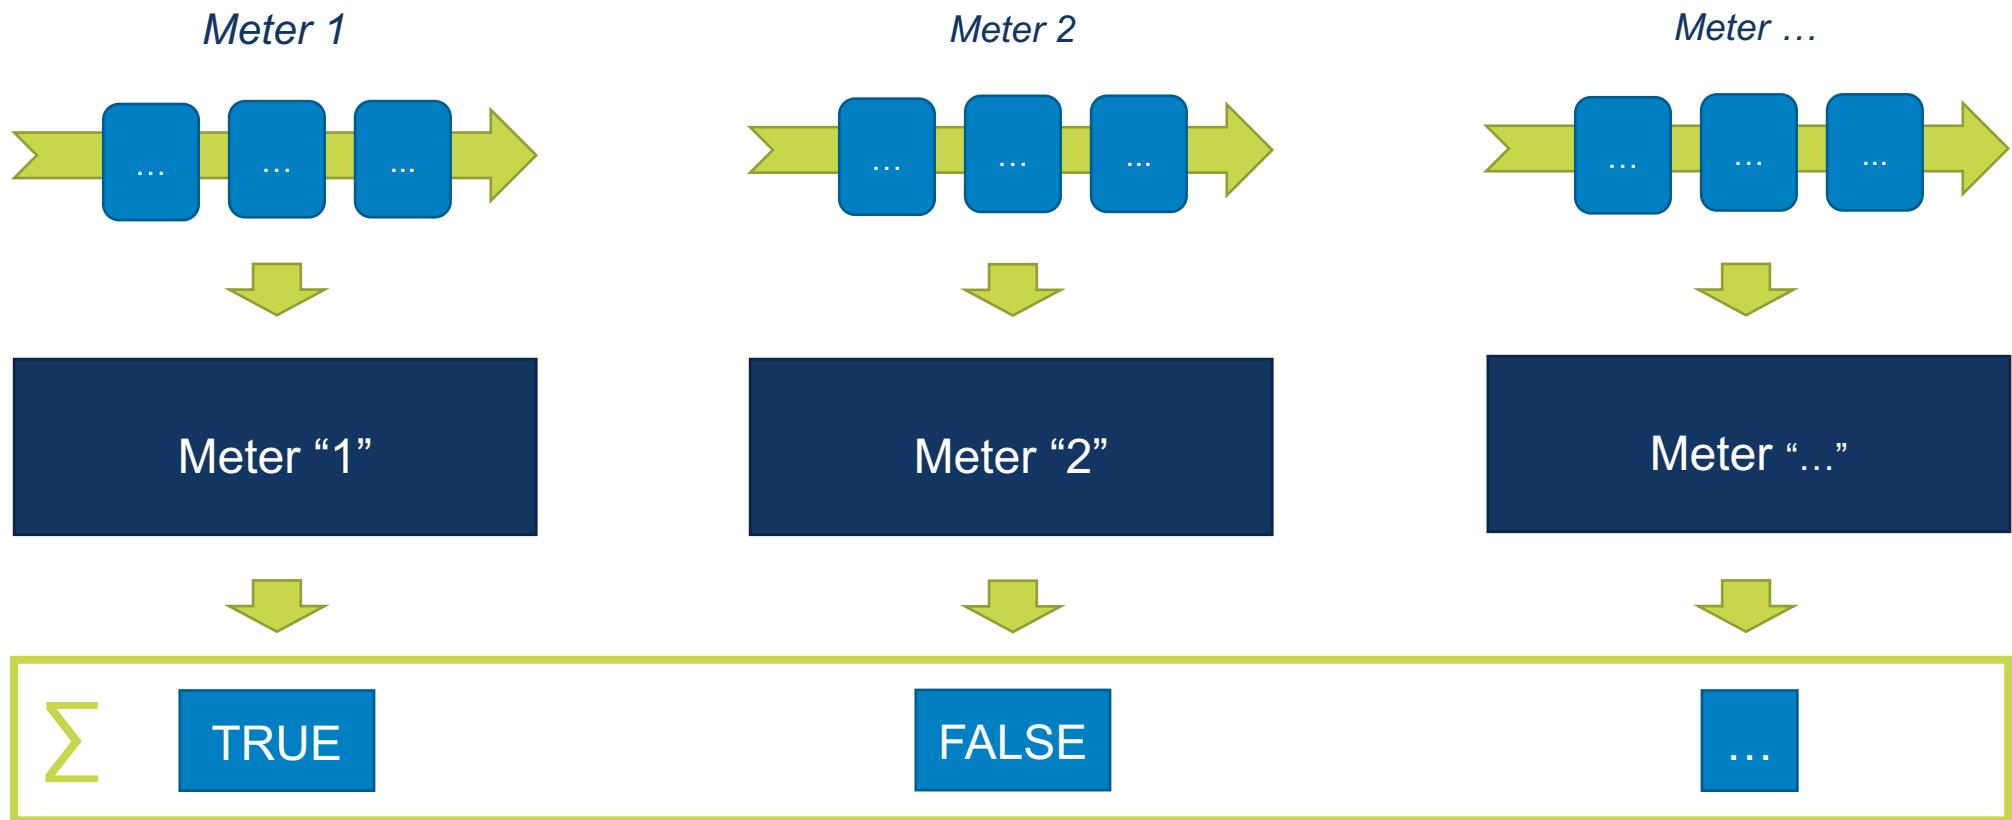
COSMOS DB



# ▶ Querying event streams



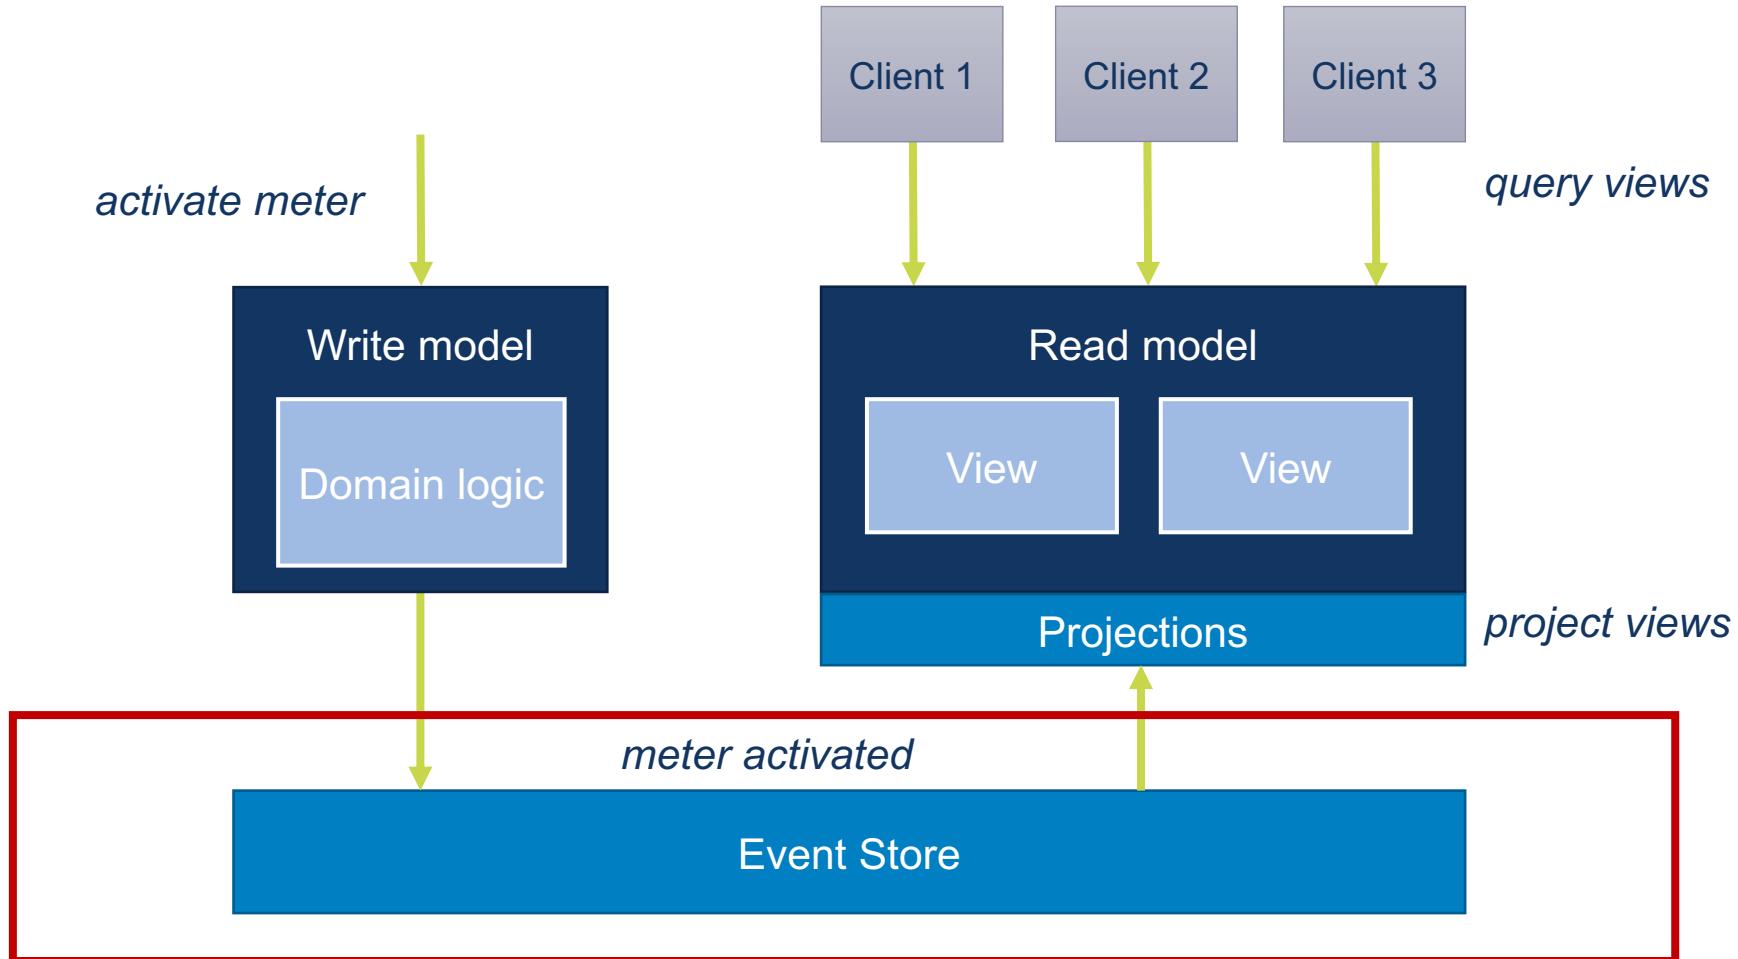
## ► Multi-stream queries



# Command Query Responsibility Segregation

Data model

# Command Query Responsibility Segregation

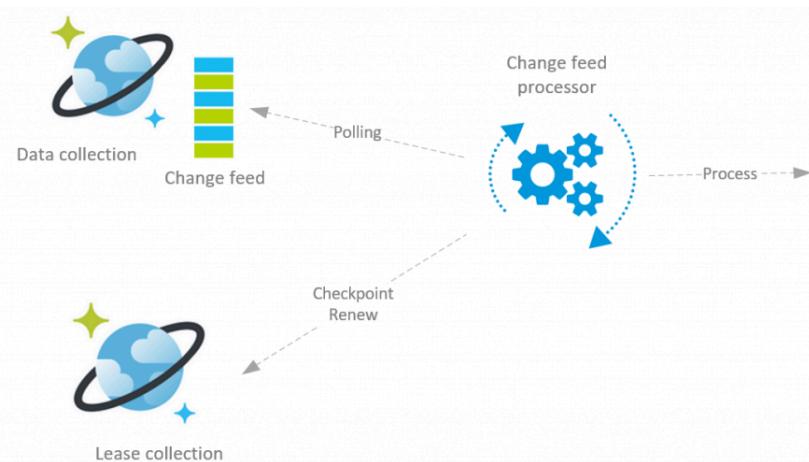


# ▲ Cosmos DB change feed

- Subscribe to changes in a Cosmos DB container
- Change feed items come in the order of their modification time
- If the data is not deleted, it will remain in the change feed



# Handle changes



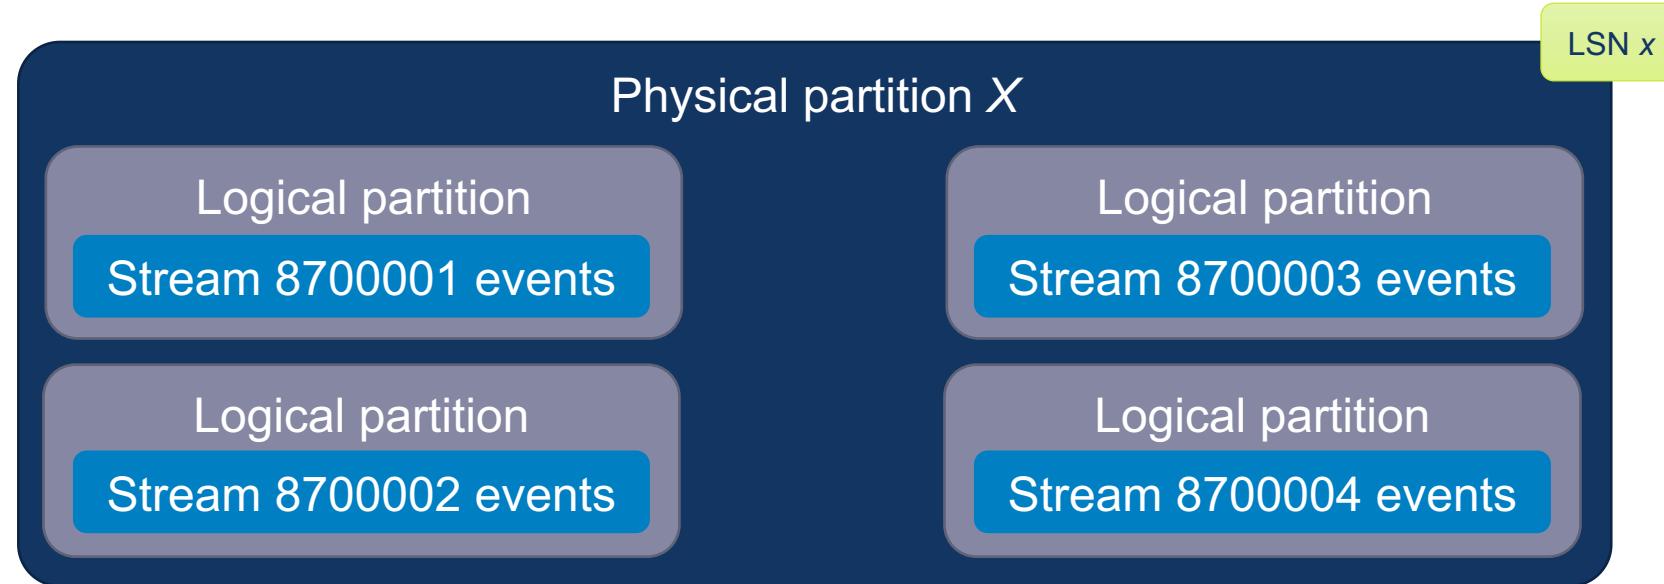
1. Deserialize events
2. For each event:
  1. Get subscribed projections
  2. For each subscribed projection:
    1. Load view
    2. Skip projection if view LSN is newer than event LSN
    3. Run projection
    4. Save view (use Etag for optimistic concurrency control)

>>

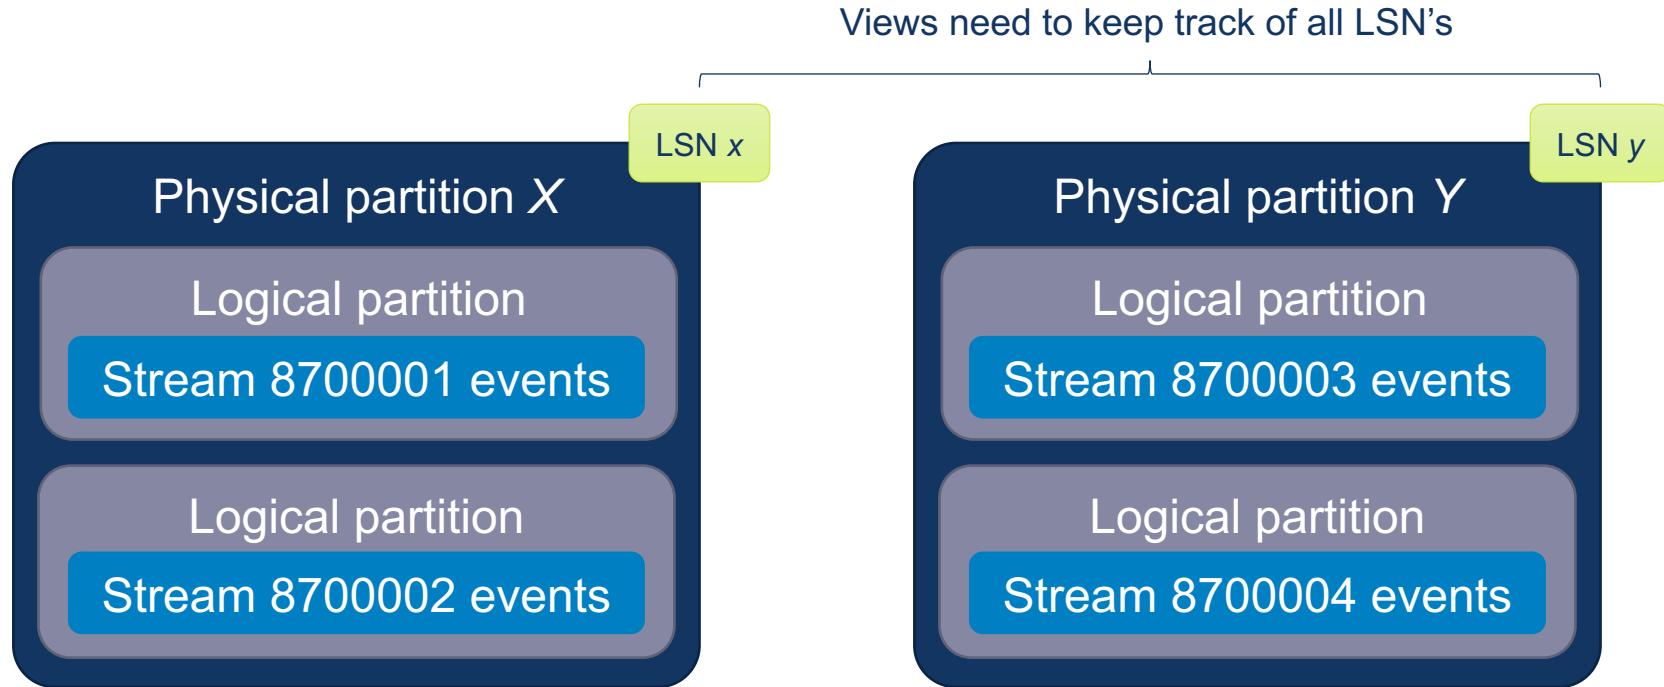
# CQRS

COSMOS DB - DEMO

## ► Cross-partition projection caveat



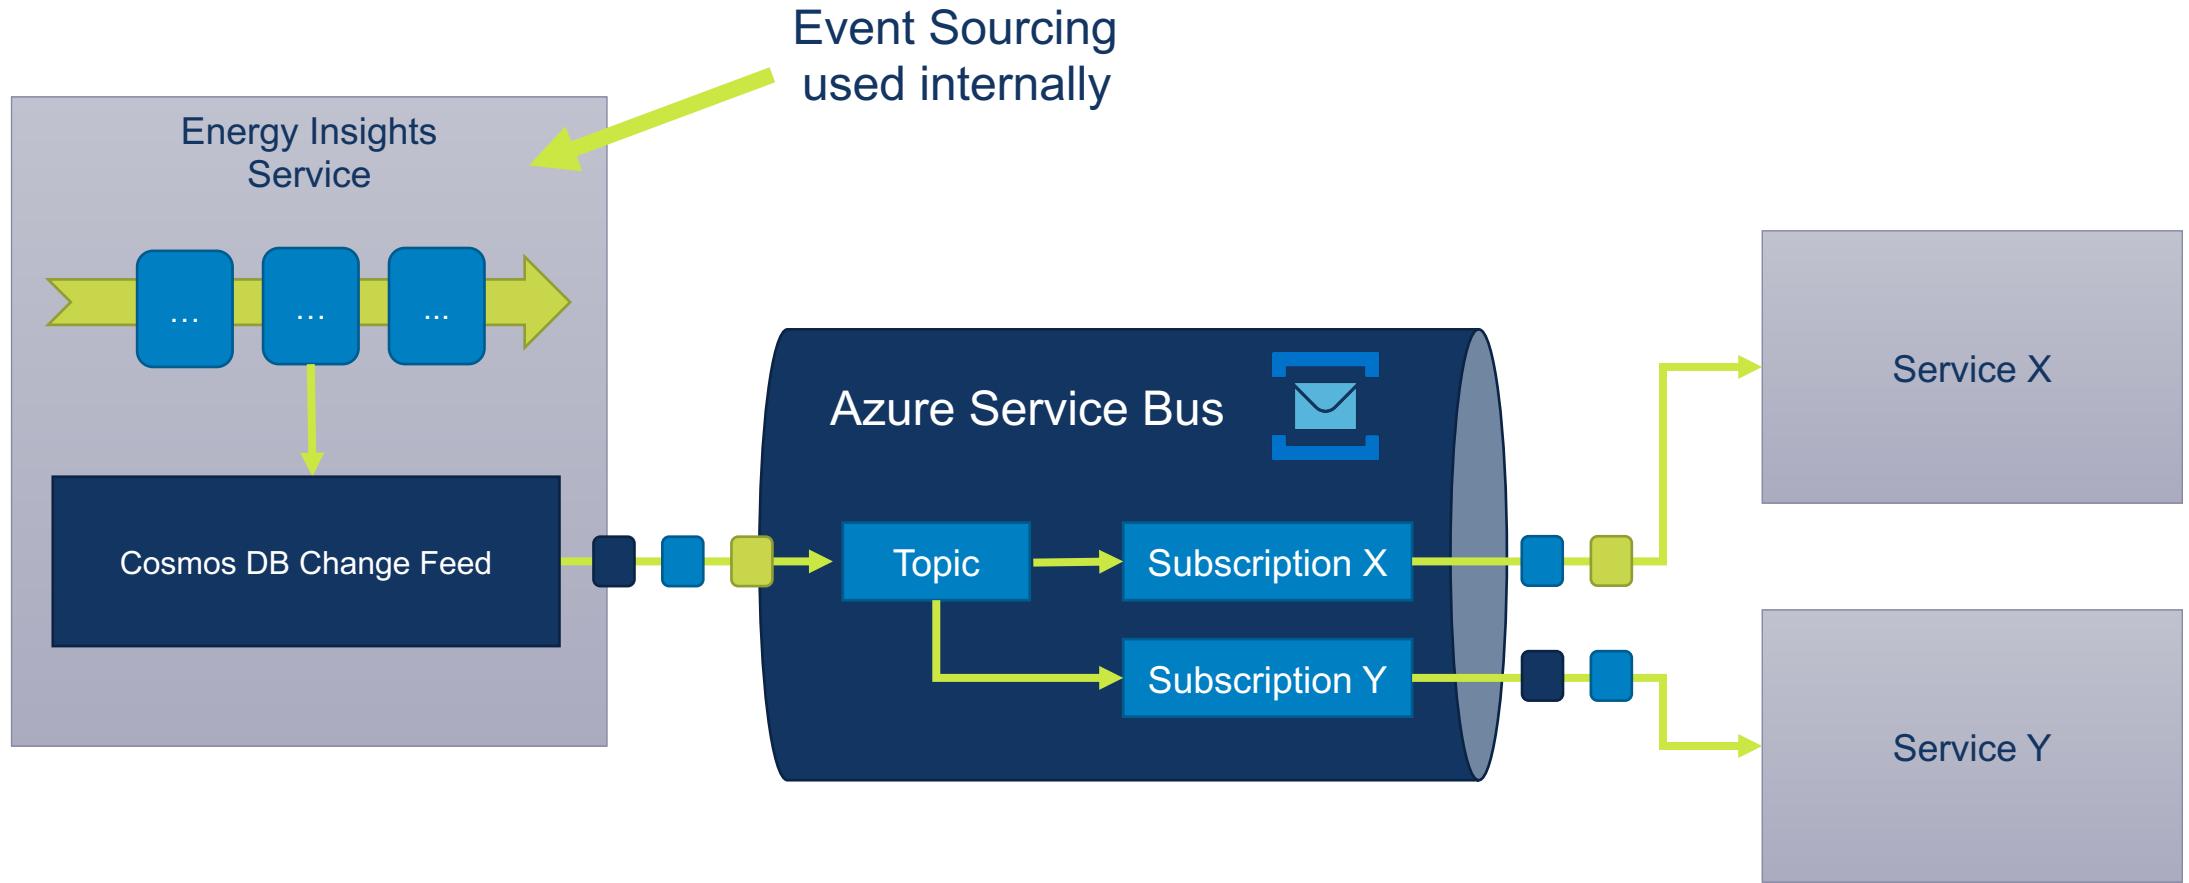
# ► Cross-partition projection caveat



V3 change feed support doesn't expose that information, possible solutions include:

- Use the V2 Change Feed Processor Library which does expose the PartitionId
- Wait for future releases to reinstate Context:  
<https://github.com/Azure/azure-cosmos-dotnet-v3/issues/1122>

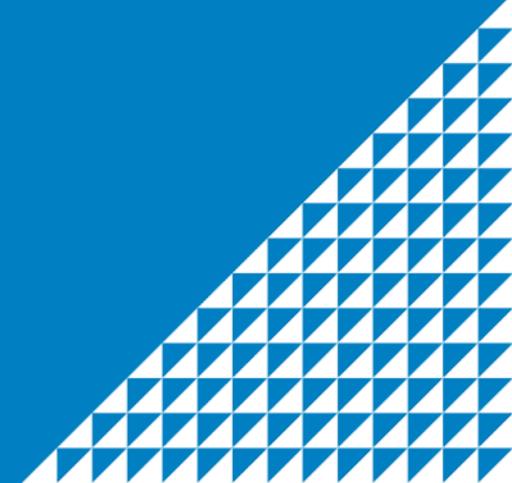
# Integration with Azure Service Bus



Event Driven Architecture to connect services

>>

# Write-side Performance



# ▲ Performance

- Partition streams
  - This happens in the real world too, e.g. closing a financial year

Meter:87000001

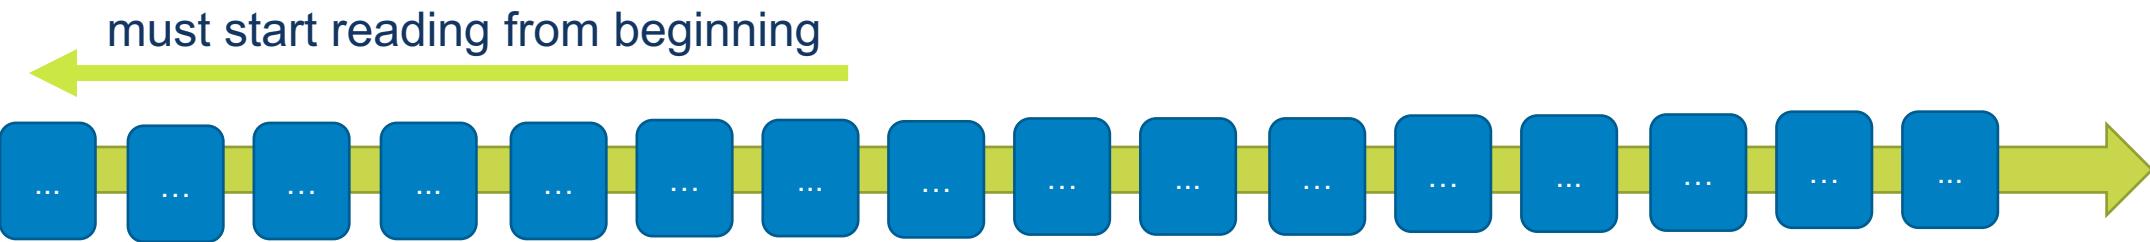


Meter:87000001:2019

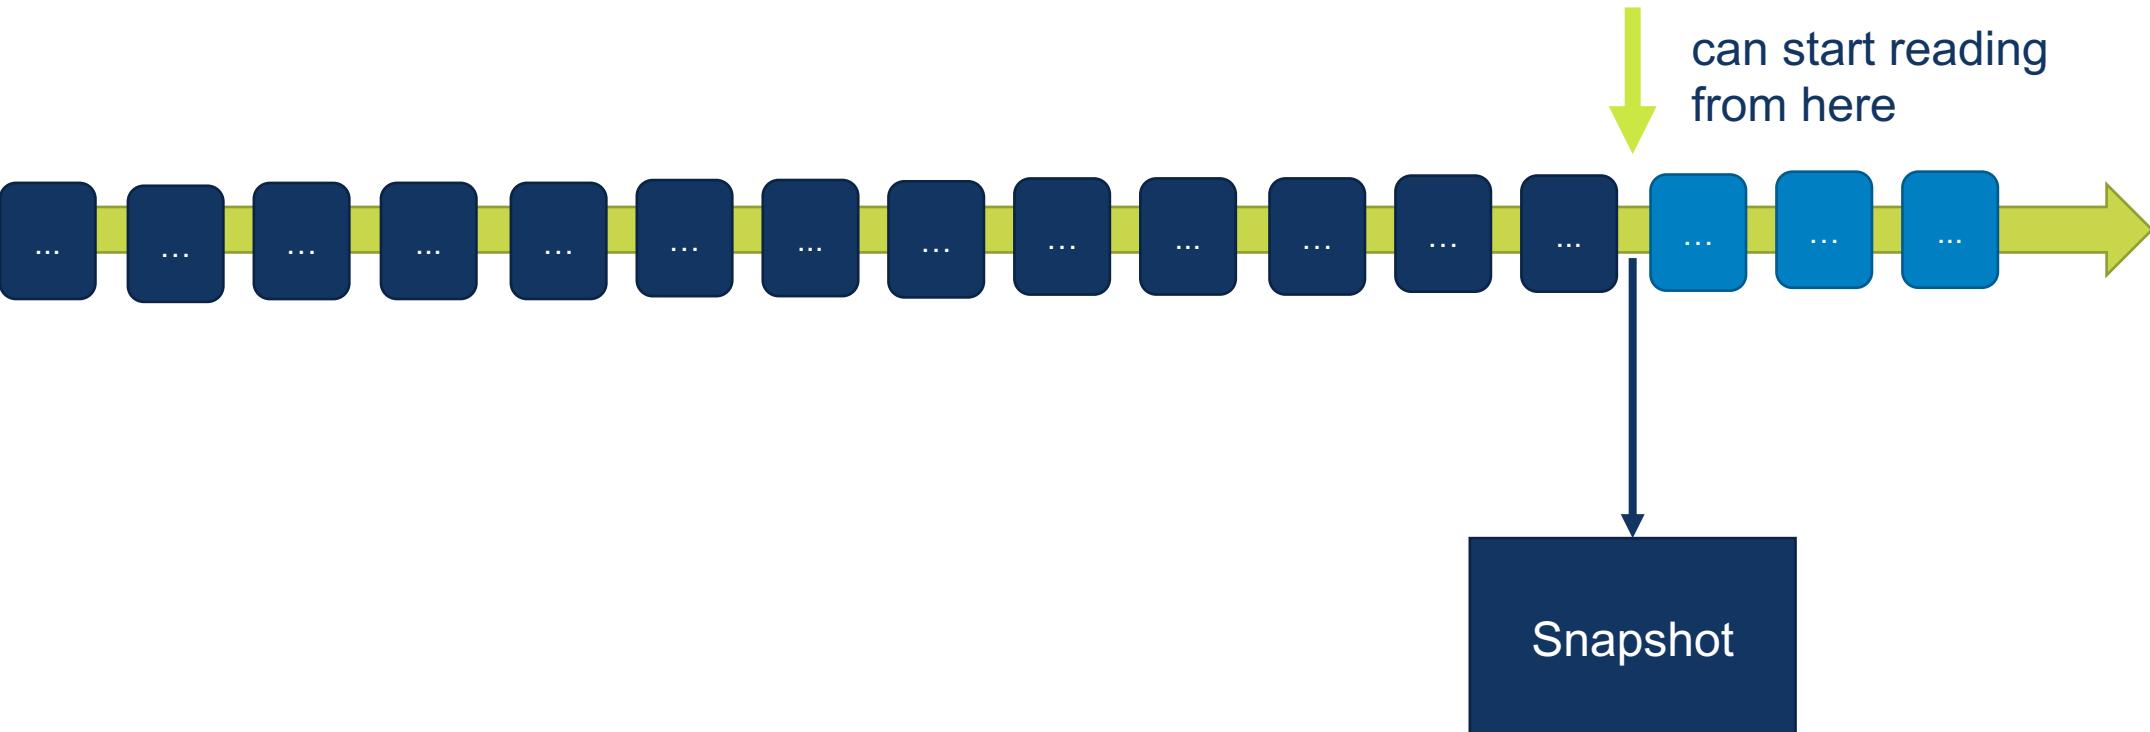
Meter:87000001:2020

- Use a cache
  - No complex cache invalidation because streams are immutable
- Use snapshots

# ► Snapshots

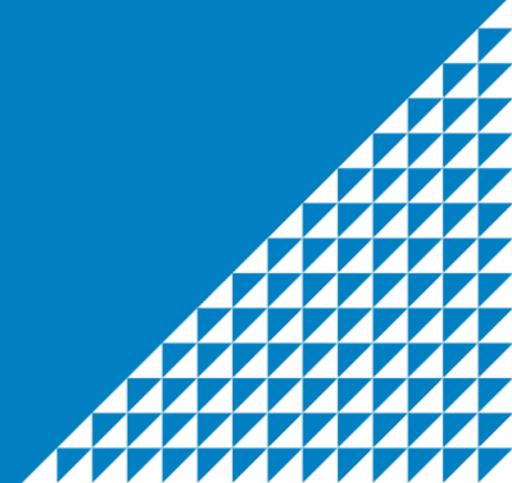


# ▶ Snapshots



>>

# GDPR



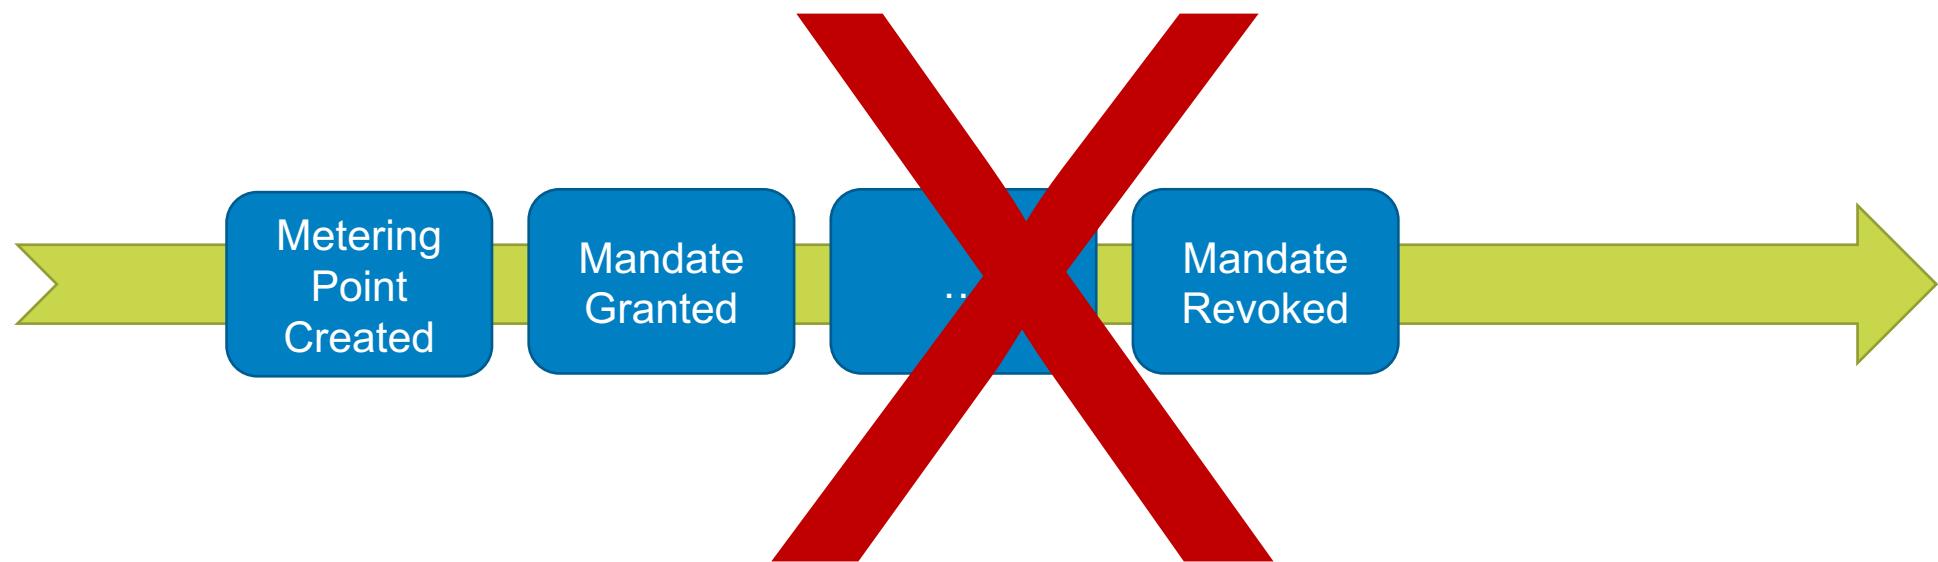
## ► Immutability vs GDPR

*“... the data subject shall have the right to obtain from the controller the erasure of personal data concerning him or her without undue delay and the controller shall have the obligation to erase personal data without undue delay...”*

Article 17 – “right to erasure”

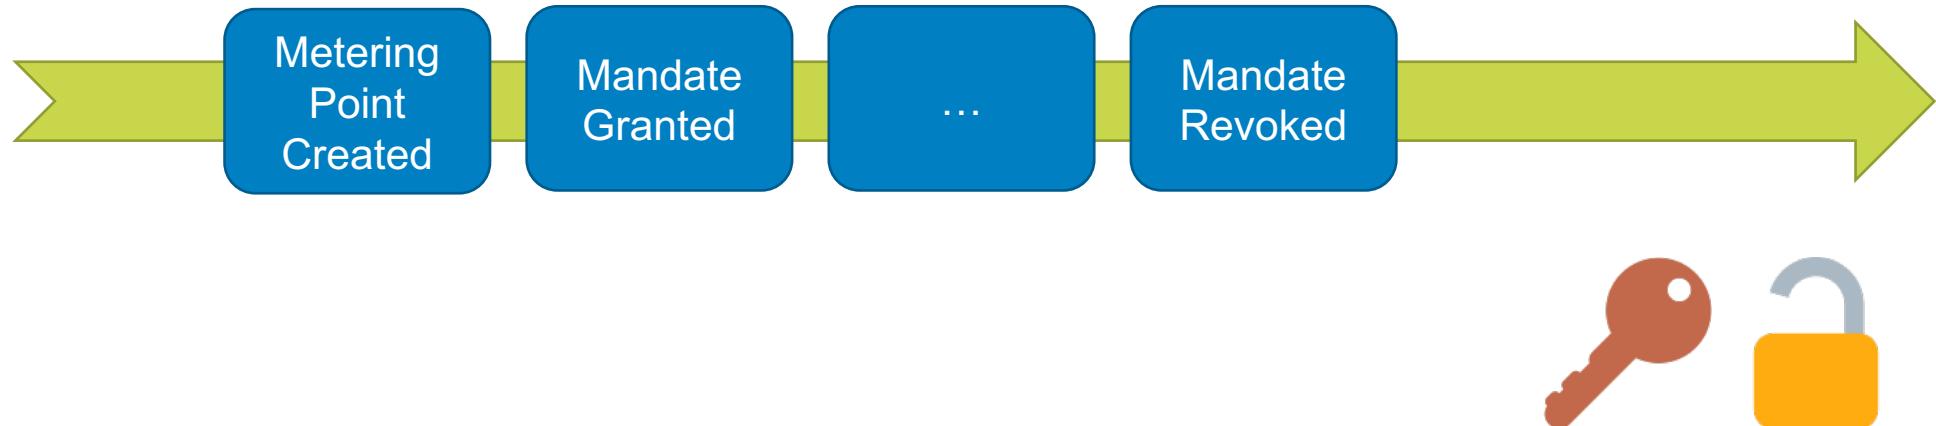
## ▲ Delete entire stream

Removing an entire stream is a “safe” operation.



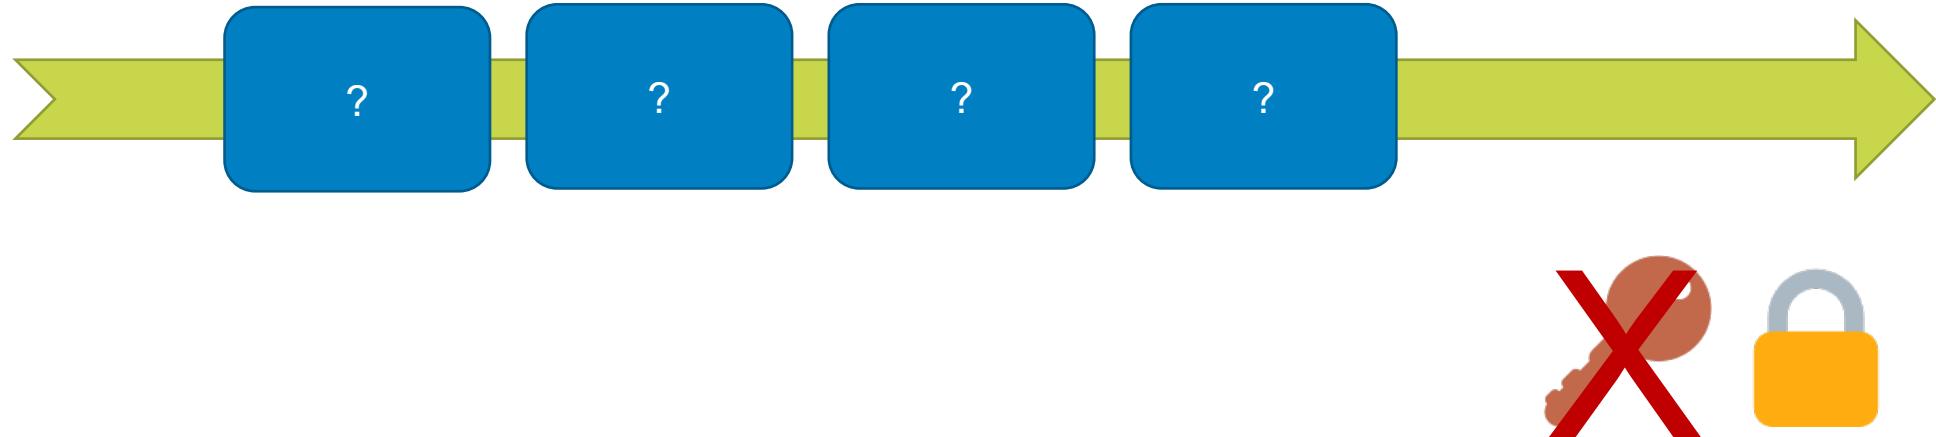
# ◀ Crypto shedding

- Encrypt each stream using a unique key.



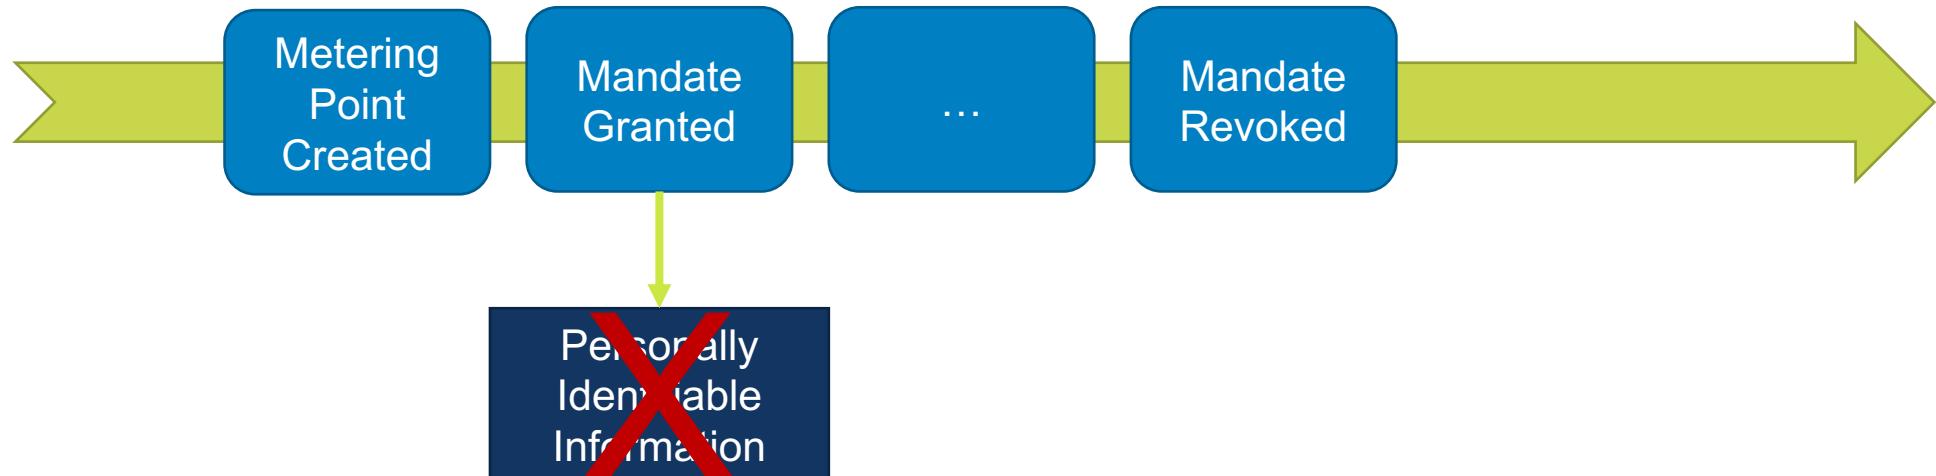
## ◀ Crypto shedding

- Encrypt each stream using a unique key.
- Throw away the key to make the PII unreadable.



## External storage

Store the PII outside of the event.



>>

All  
usernames/e-mail addresses/whatever  
must be unique!

## ► All usernames/e-mail addresses/whatever must be unique!

- Don't feel obligated to use ES everywhere in your application!
- Use the unique field as the event stream identifier
  - We use the unique meter id as the stream id
- Keep an index that you can reference during command processing
  - Only use on command side, do not mix with read model to avoid unnecessary dependencies
- What's the real business scenario?
  - How often does this really happen? Can you just do a compensating action?

# Thank you!



Sander Molenkamp

 @amolenk

All demo source code available at:

<https://github.com/amolenk/CosmosEventSourcing>



dotnetFlix